

Graphs

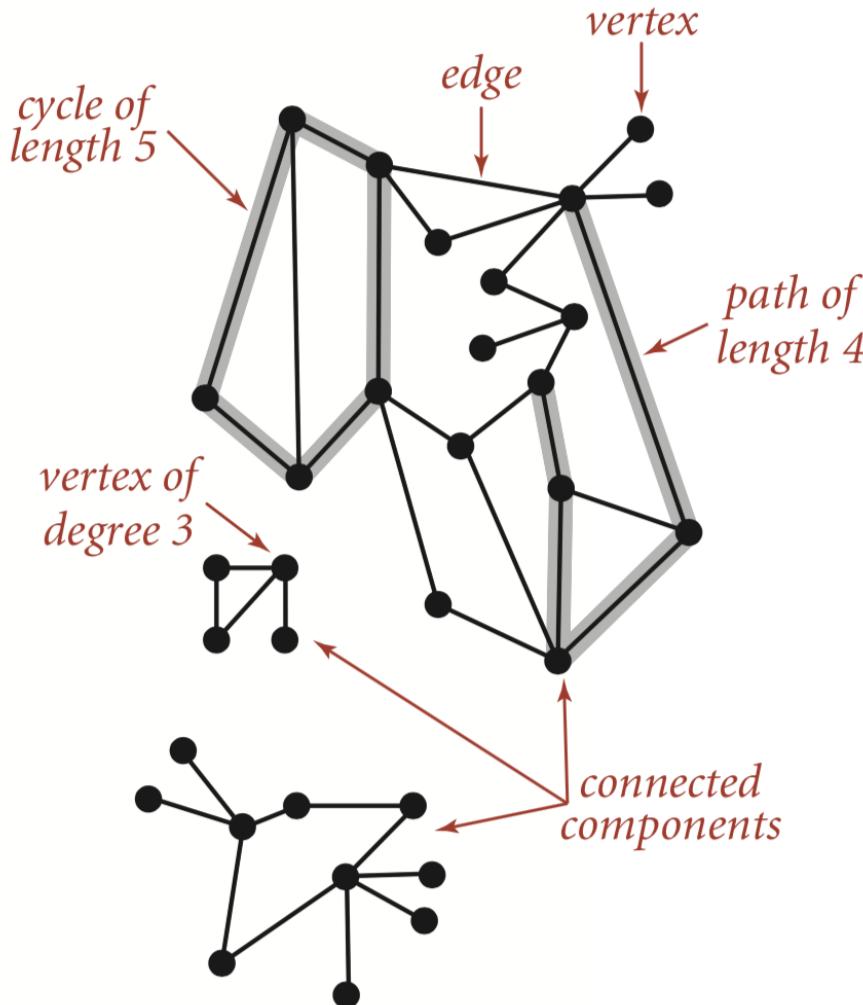
It's all about relationships

Mustafa Hajij
MSDS program
University of San Francisco

What's a graph?

- A graph is a collection of connected element pairs, $u \rightarrow v$ or $u - v$
- As with a tree, a graph is the aggregate of nodes/edges
- Nodes can be email addresses, map locations, documents, tasks to perform, URLs on the web, customers, computers on network, friends, observations, sensors, states in markov chain, ...
- Terms: *nodes* or *vertices* connected with *edges*, which can have labels; e.g., recall the Trie graph with labeled edges
- *Directed* graphs have arrows as edges but *undirected* use lines
- For n nodes, num directed edges is ≥ 0 and $\leq n^2$ since n nodes could connect to all other nodes and itself; if no self cycles then the max is $n(n - 1)$

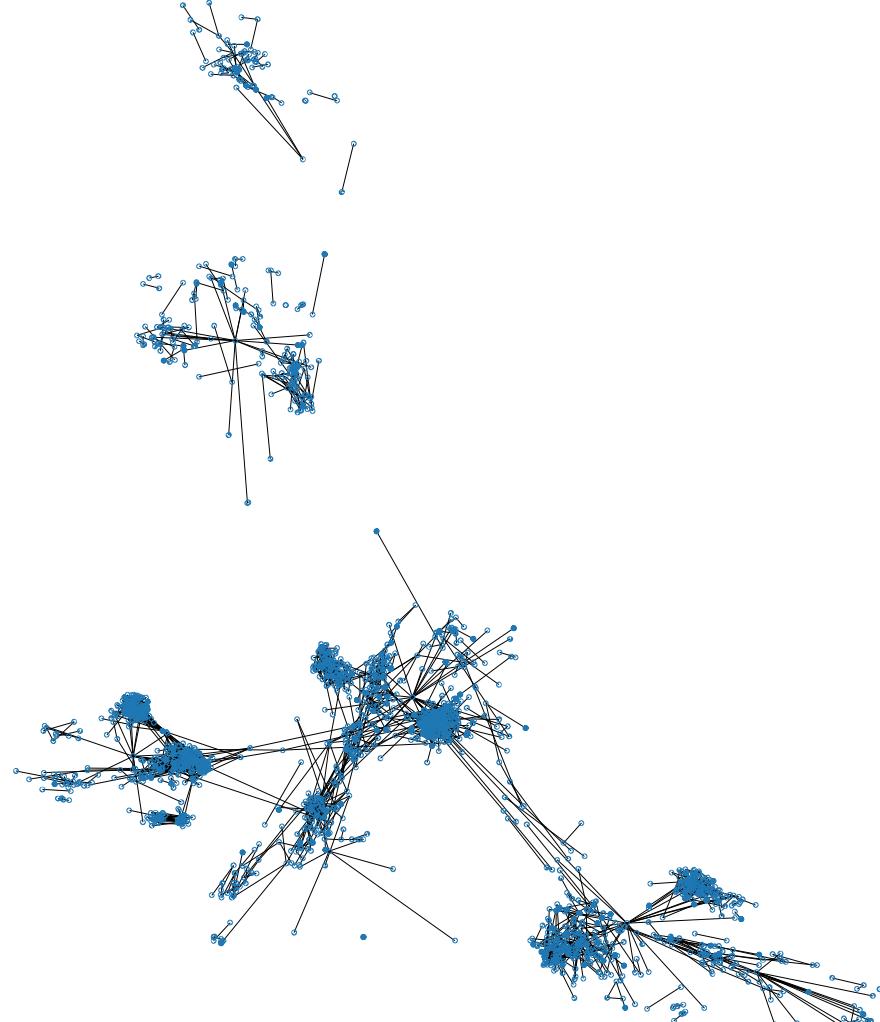
Undirected graph, terms



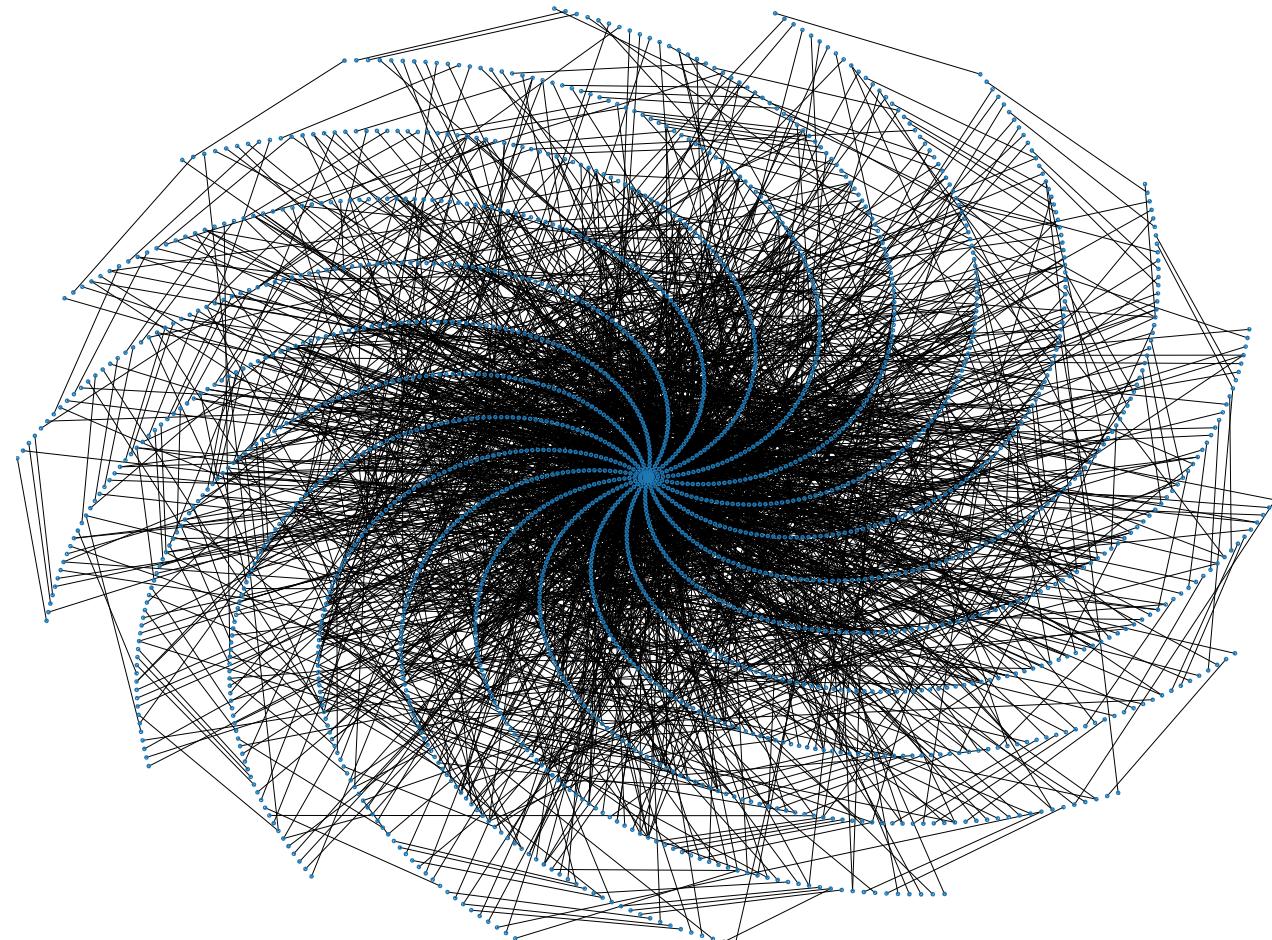
From: Algorithms book by Robert Sedgewick and Kevin Wayne

Facebook friend network, different layouts

Fruchterman Reingold



spiral



UNIVERSITY OF SAN FRANCISCO

Common questions

- Is q reachable from p?
- How many edges are on paths between p and q?
- Is graph connected? (reach any p from any q)
- Is graph cyclic? (p reaches p traversing at least one edge)
- Which nodes are within k edges of node p? (neighborhood)
- What is shortest path (num edges) from p to q?
- What is shortest path using edge weights? [beyond scope of 689]
- Traveling salesman problem [beyond scope of 689]
https://en.wikipedia.org/wiki/Travelling_salesman_problem

Adjacency matrix implementations

- Adjacency matrix, $n \times n$ matrix of $\{0,1\}$ if unlabeled or $\{\text{labels}\}$ if edges are labeled; undirected matrices are symmetric
- Wastes space for sparse edges; use sparse matrix
- Fast to access arbitrary node's edges if we have unique node IDs

	a	b	c	d	e	f	g	h	i	j	k	l	m
a	0	1	0	0	1	0	0	0	0	0	0	0	0
b	1	0	1	0	1	1	0	0	0	0	0	0	0
c	0	1	0	1	0	1	1	0	0	0	0	0	0
d	0	0	1	0	0	0	1	0	0	0	0	0	0
e	1	1	0	0	0	1	1	1	0	0	0	0	0
f	0	1	1	0	1	0	1	0	0	0	0	0	0
g	0	0	1	1	1	1	0	0	1	0	0	0	0
h	0	0	0	0	1	0	0	0	0	0	0	0	0
i	0	0	0	0	0	0	1	0	0	0	0	0	0
j	0	0	0	0	0	0	0	0	0	1	1	1	1
k	0	0	0	0	0	0	0	0	1	0	1	0	0
l	0	0	0	0	0	0	0	0	1	1	0	1	0
m	0	0	0	0	0	0	0	0	0	1	0	1	0

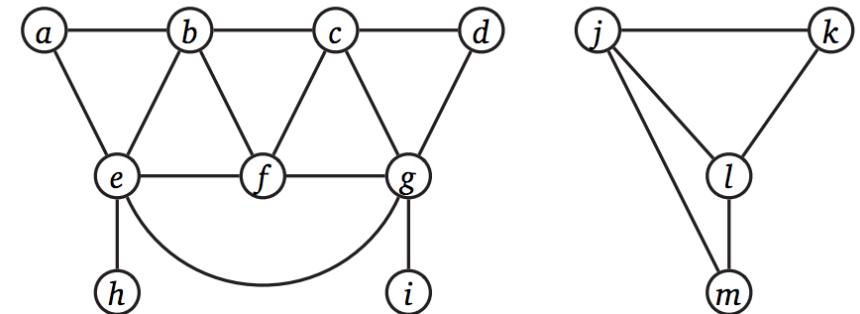


Figure 5.11. An adjacency matrix for our example graph.

Adjacency list implementations

- List of edge lists for nodes
- Fast arbitrary node access for numbered nodes, space efficient

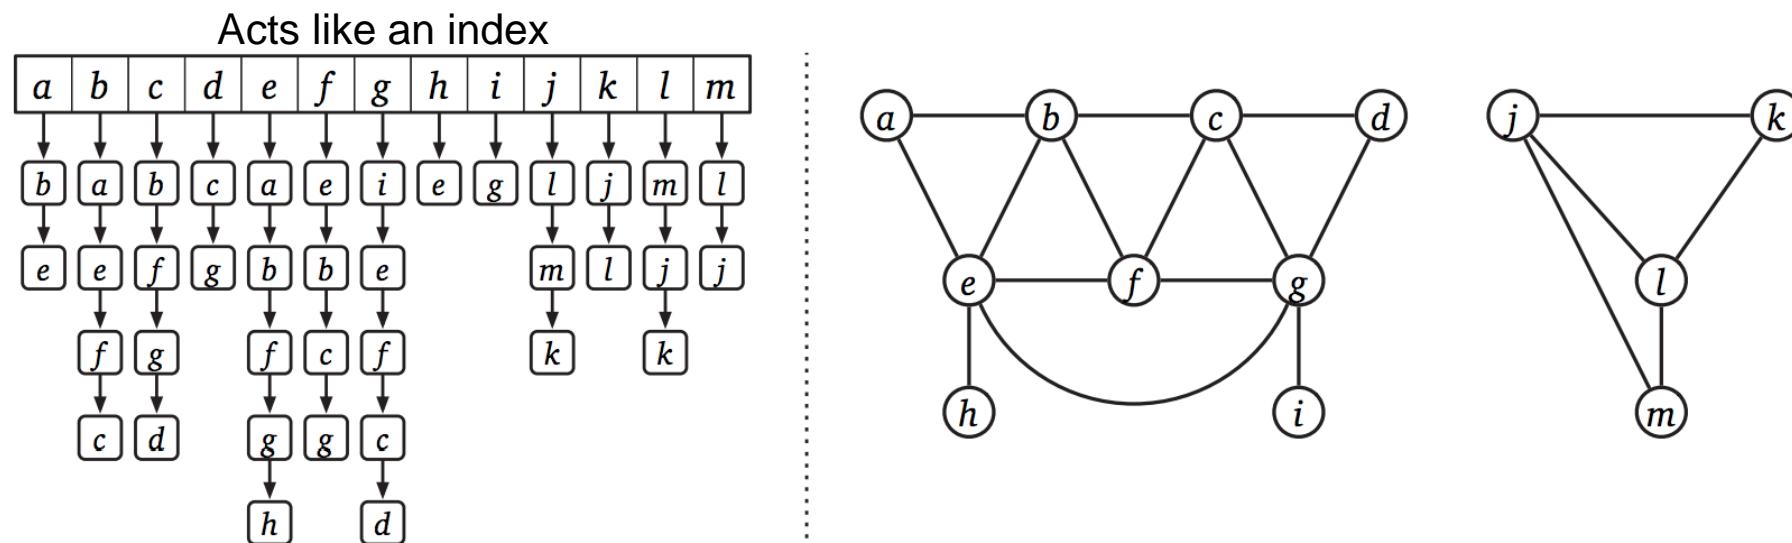
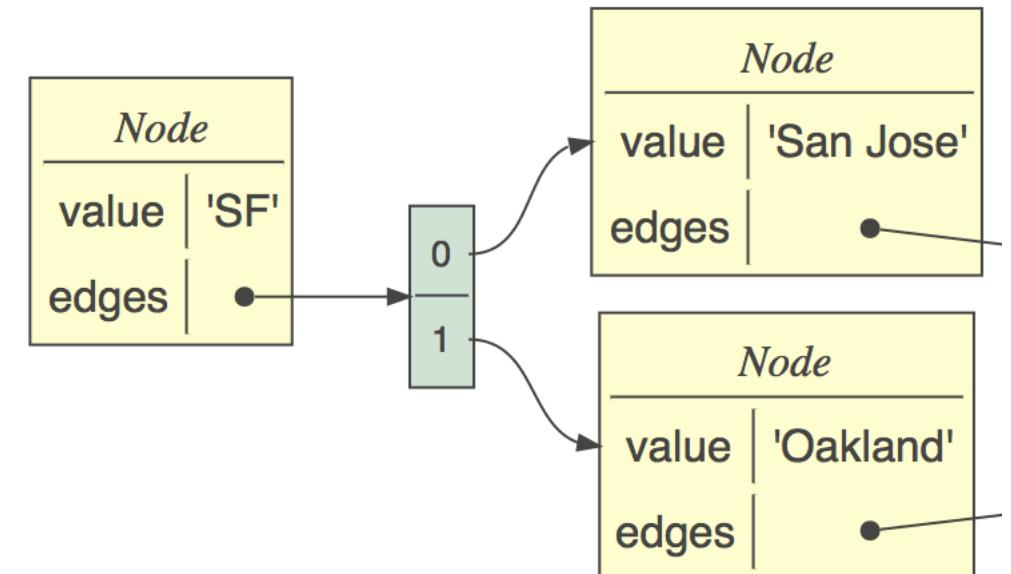


Figure 5.9. An adjacency list for our example graph.

Connected nodes implementation

- Common implementation due to nice mapping to objects
- Each node has info about that node and its edge list
- Use list or dictionary index if you need to access nodes directly

```
class Node:  
    def __init__(self, value):  
        self.value = value  
        self.edges = []  
    def edge(self, target:Node):  
        self.edges.append(target)
```

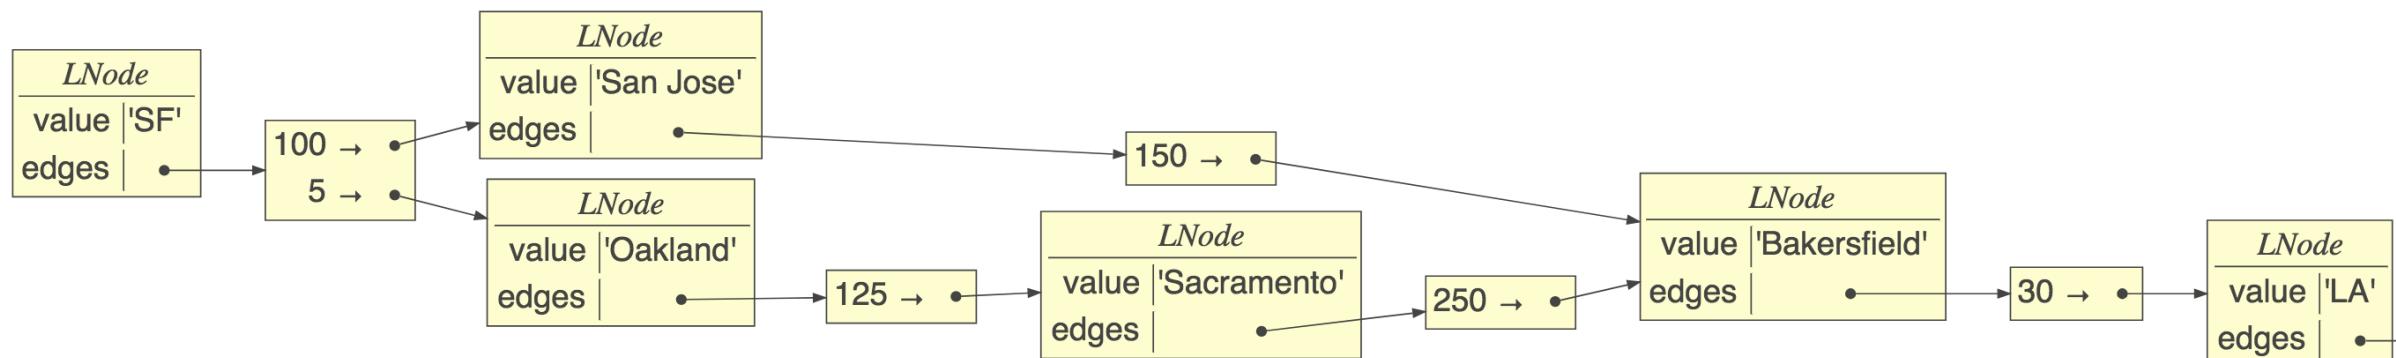


Implementation with labels

```
class LNode:  
    def __init__(self, value):  
        self.value = value  
        self.edges = {}  
    def edge(self, label, target):  
        self.edges[label] = target
```

Edge->node is a dictionary, not list

sf.edge(100,sj)
sj.edge(150,baker)
...

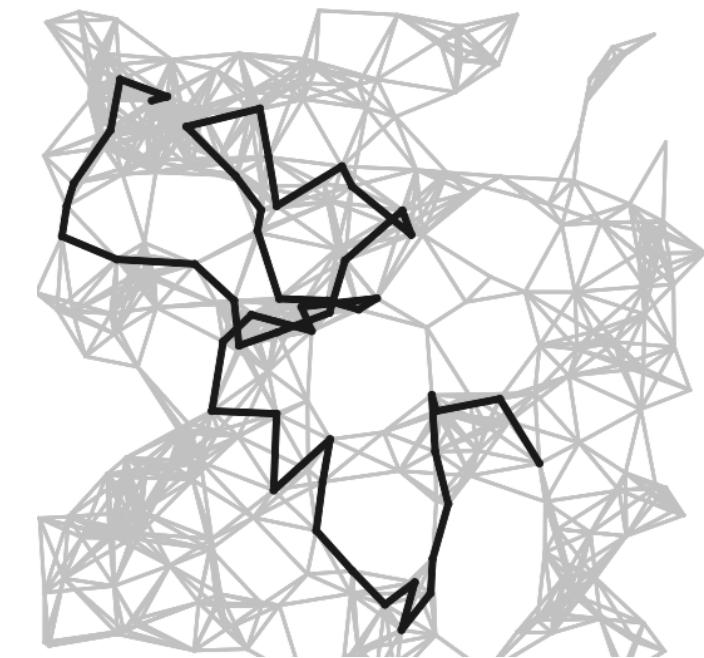


Depth-first search (review)

- The fundamental algorithm for answering graph questions
- Visits all reachable nodes from p, avoiding cycles
- Go deep first

```
def walk_graph(p:Node, visited=set()):  
    if p in visited: return  
    visited.add(p)  
    for q in p.edges:  
        walk_graph(q, visited)
```

```
RECURSIVEDFS(v):  
    if v is unmarked  
        mark v  
        for each edge vw  
            RECURSIVEDFS(w)
```



$O(n,m) = n + m$, for n **nodes**, m **edges**; m can be n^2
(counts node visits, edge traversals)

Is there a cycle (p can reach p)?

- E.g., can I get back to starting spot via one-way streets?
- If we run into starting node in visited set, return True;
blue means different than plain walk:

```
def iscyclic(p:Node) -> bool:  
    return iscyclic_(p,p,set())
```

```
def iscyclic_(start:Node, p:Node, visited) -> bool:  
    if p in visited:  
        if p is start: return True # we find start?  
        return False # can't loop forever so don't chase  
    visited.add(p)  
    for q in p.edges:  
        c = iscyclic_(start, q, visited)  
        if c: return True # found it, we can stop  
    return False
```

Find set of nodes p can reach

- E.g., can we reach city y from city x (hint: Juneau not accessible)?
- Need two sets, one for avoiding cycles, another for reached nodes
- If we used visited for both, then p would also appear to reach itself, which might not be true; we add start node to visited but doesn't mean we can return to start via any path in graph

```
def reachable(p:Node) -> set:  
    reaches = set();  
    reachable_(p, reaches, set())  
    return reaches  
  
def reachable_(p:Node, reaches:set, visited:set):  
    if p in visited: return  
    visited.add(p)  
    for q in p.edges:  
        reaches.add(q) # add only if we traverse  
        reachable_(q, reaches, visited)
```

Find set of nodes p can reach, track depth

- E.g., how many hops from person A to B in social network?
- Track node->depth map, not just set of nodes

```
def depths(p:Node) -> dict:  
    reaches = {}  
    depths_(p, reaches, depth=0)  
    return reaches  
  
def depths_(p:Node, reaches:dict, depth:int):  
    if p in reaches: return  
    reaches[p] = depth # distance to p from start (could be 0)  
    for q in p.edges:  
        depths_(q, reaches, depth+1)
```

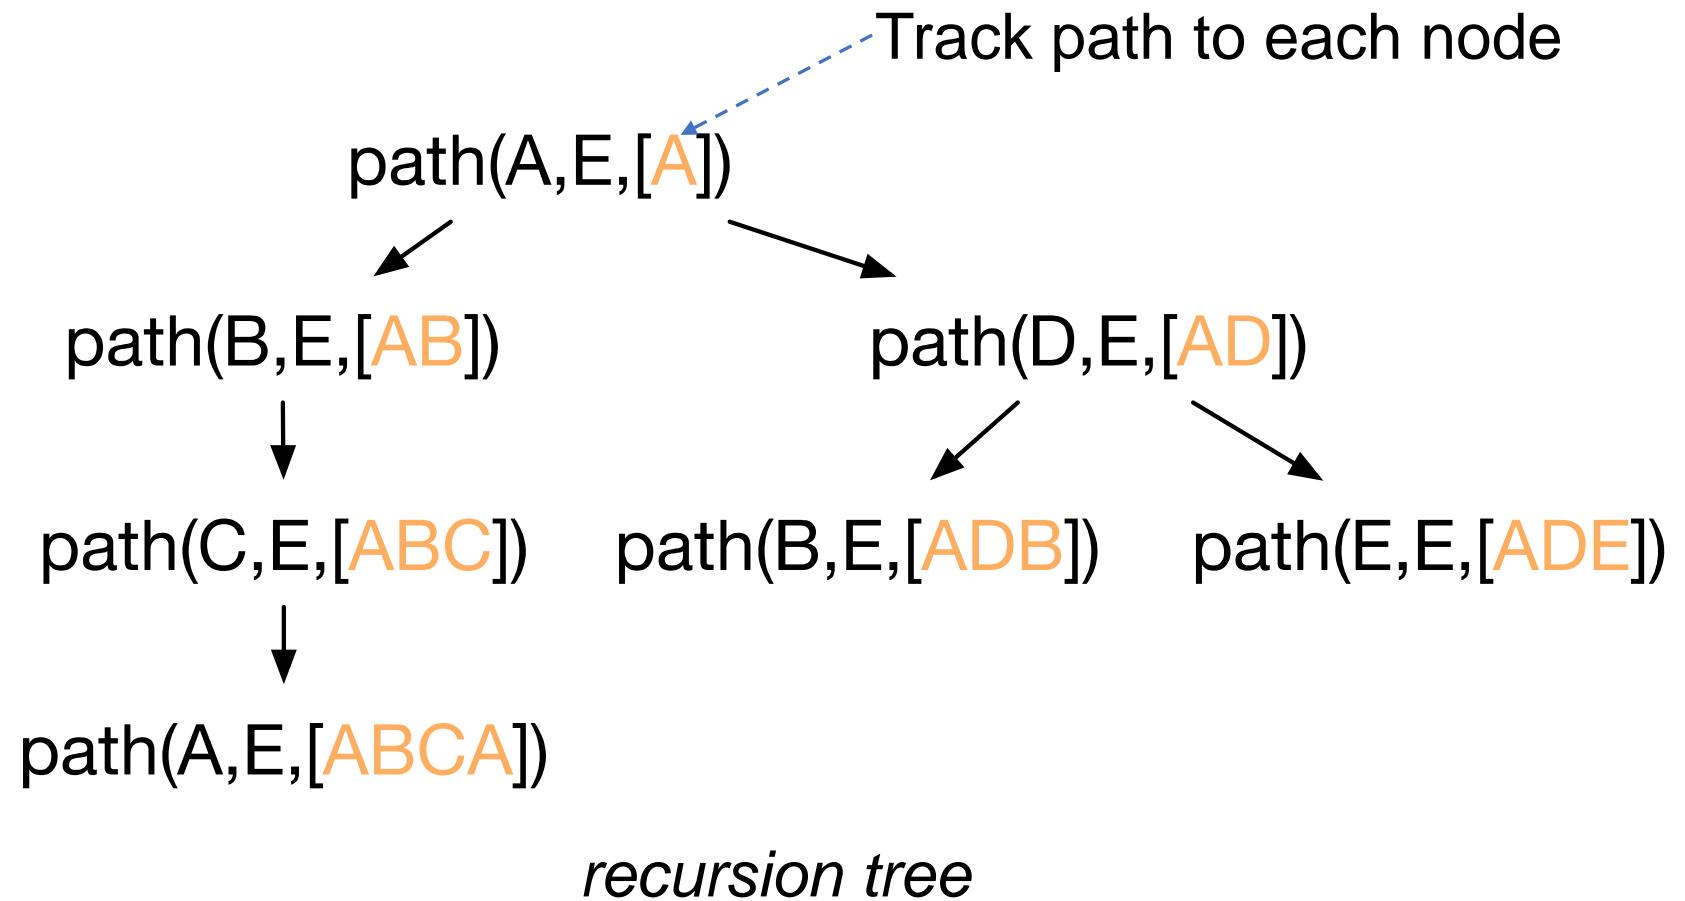
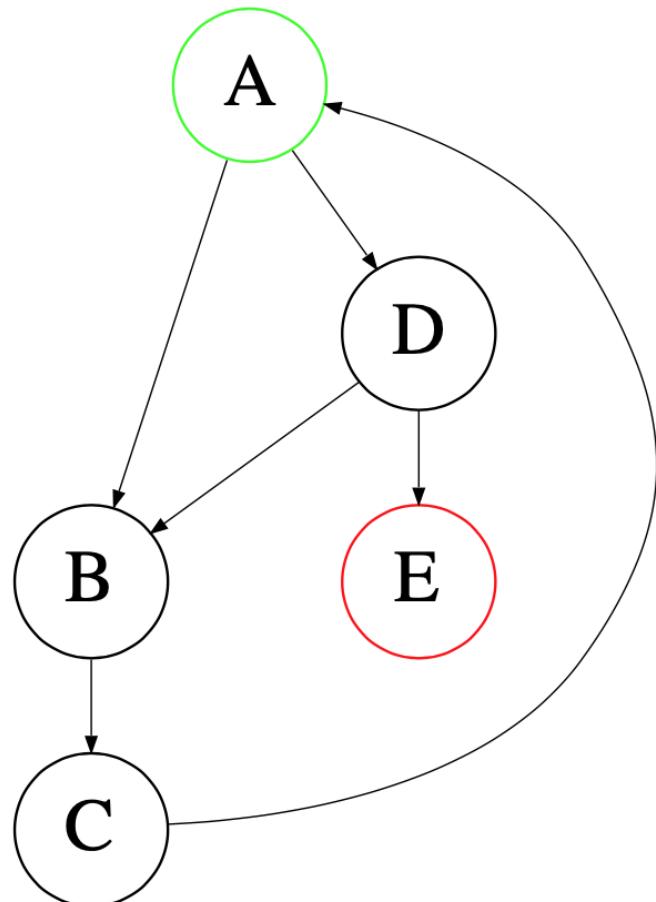
Find neighborhood within k edges

- E.g., find all friends within k friend hops
- Track dict node->depth, stop when we reach depth

```
def neighbors(p:Node, k:int) -> dict:  
    reaches = dict()  
    neighbors_(p, k, reaches, depth=0)  
    return reaches
```

```
def neighbors_(p:Node, k:int, reaches:dict, depth:int):  
    if p in reaches or depth>k: return # terminate at depth  
    reaches[p] = depth  
    for q in p.edges:  
        neighbors_(q, k, reaches, depth+1)
```

Find first path from p to q

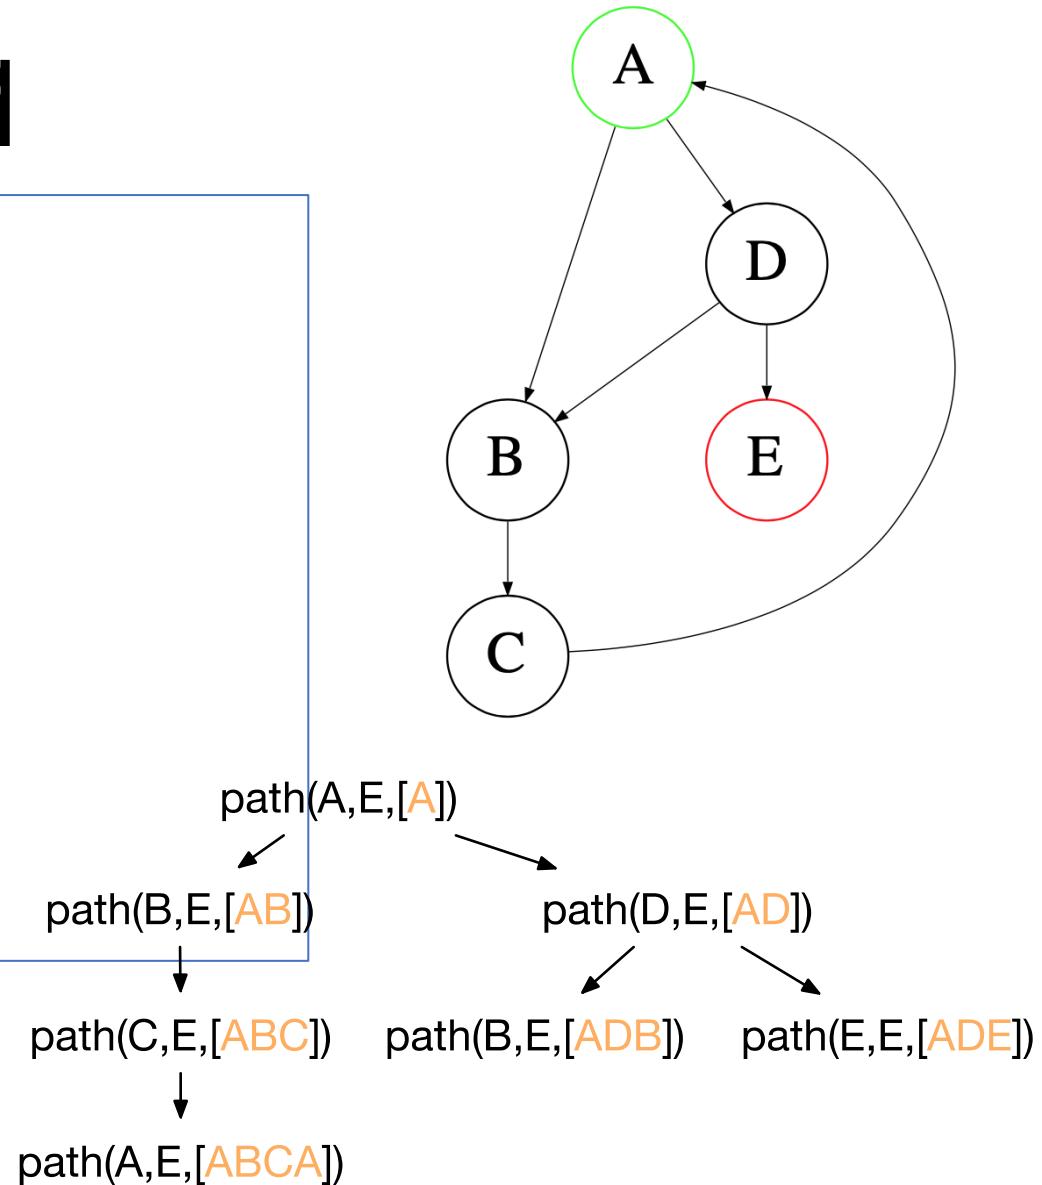


Find first path from p to q

```
def path(p:Node, q:Node) -> list:  
    return path_(p, q, [p], set())
```

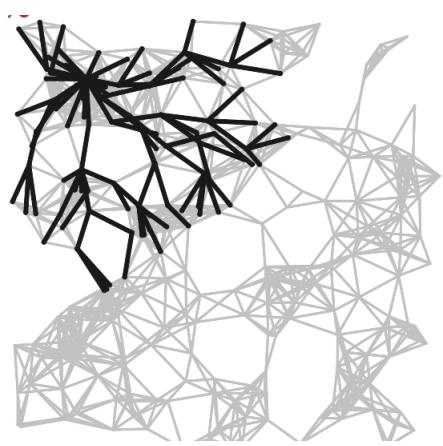
```
def path_(p:Node, q:Node, path:list, visited:set):  
    if p is q: return path      # found q!  
    if p in visited: return None # avoid cycles  
    visited.add(p)  
    for e in p.edges:  
        pa = path_(e, q, path+[e], visited)  
        if pa is not None: return pa  
    return None
```

Must track path not just set of reachable nodes

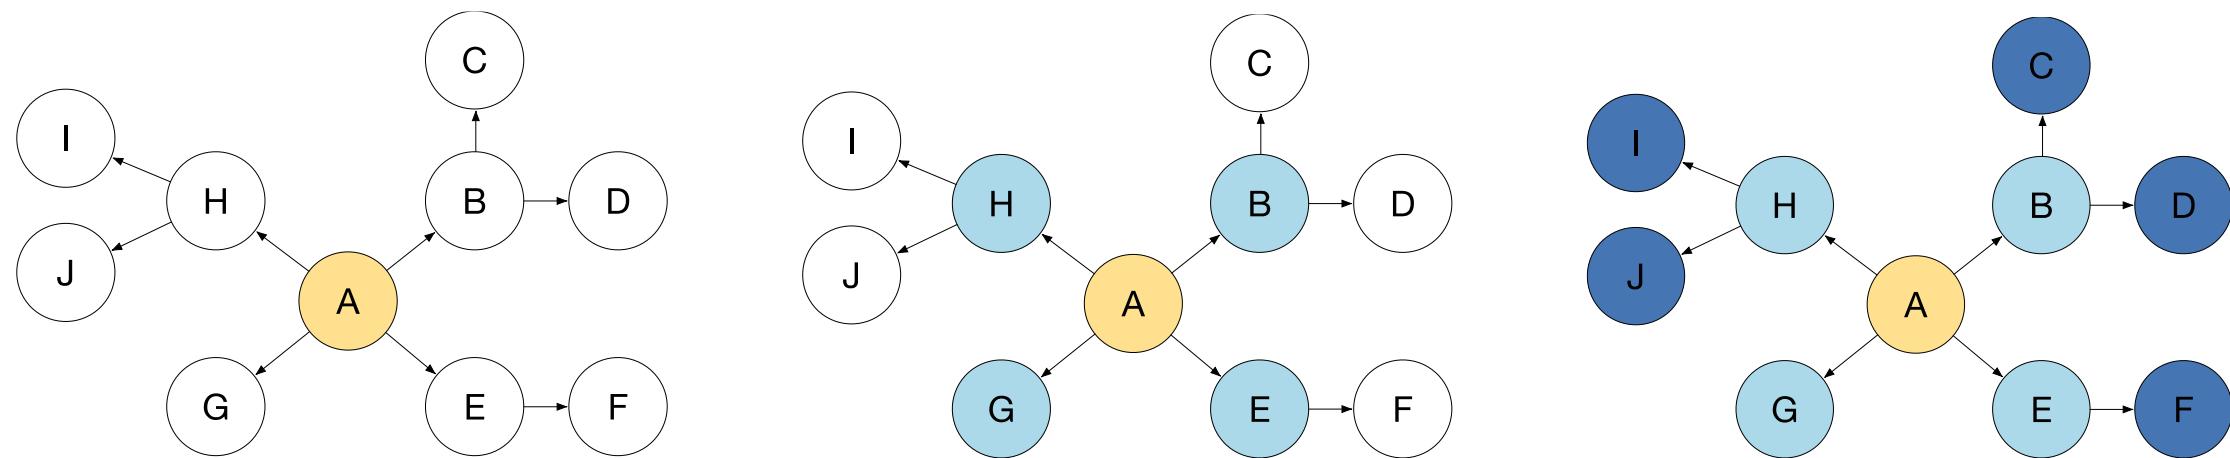


Breadth-first search vs DFS

- Visit all children then grandchildren...



Algorithms book by Sedgewick, Wayne



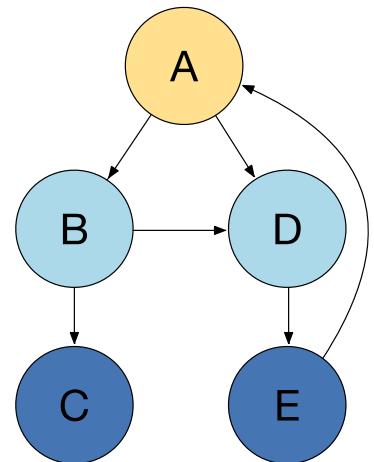
Todo: show worklist growing/shrinking here or next slide (animation)

BFS implementation

- Maintains work list of nodes and visited set
- **BFS** **DFS**

Visit A	Visit A
Visit B	Visit B
Visit D	Visit C
Visit C	Visit D
Visit E	Visit E

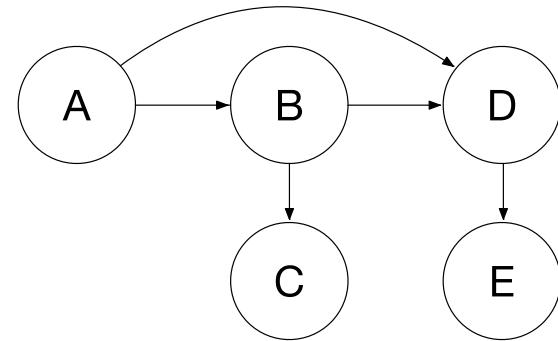
- Add to work list end, pull from front (queue)
- *Nonrecursive*



```
def BFS(root:LNode):  
    visited = {root}  
    worklist = [root]  
    while len(worklist)>0:  
        p = worklist.pop(0) # dequeue  
        print(f"Visit {p}")  
        for q in p.edges:  
            if q not in visited:  
                worklist.append(q)  
                visited.add(q)
```

Find shortest path from p to q?

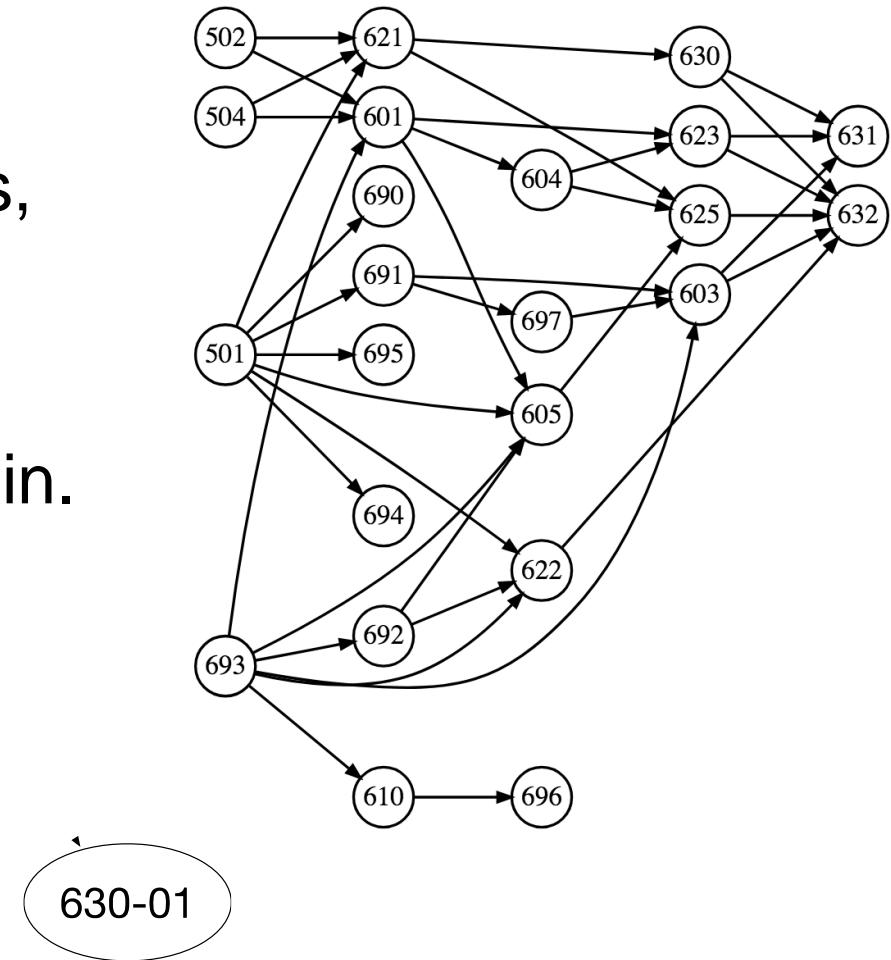
- BFS where work list is a list of paths not list of nodes
- Tail of path is where we left off working on it
- By searching all children before going deeper, we automatically find paths with shortest (unweighted) lengths



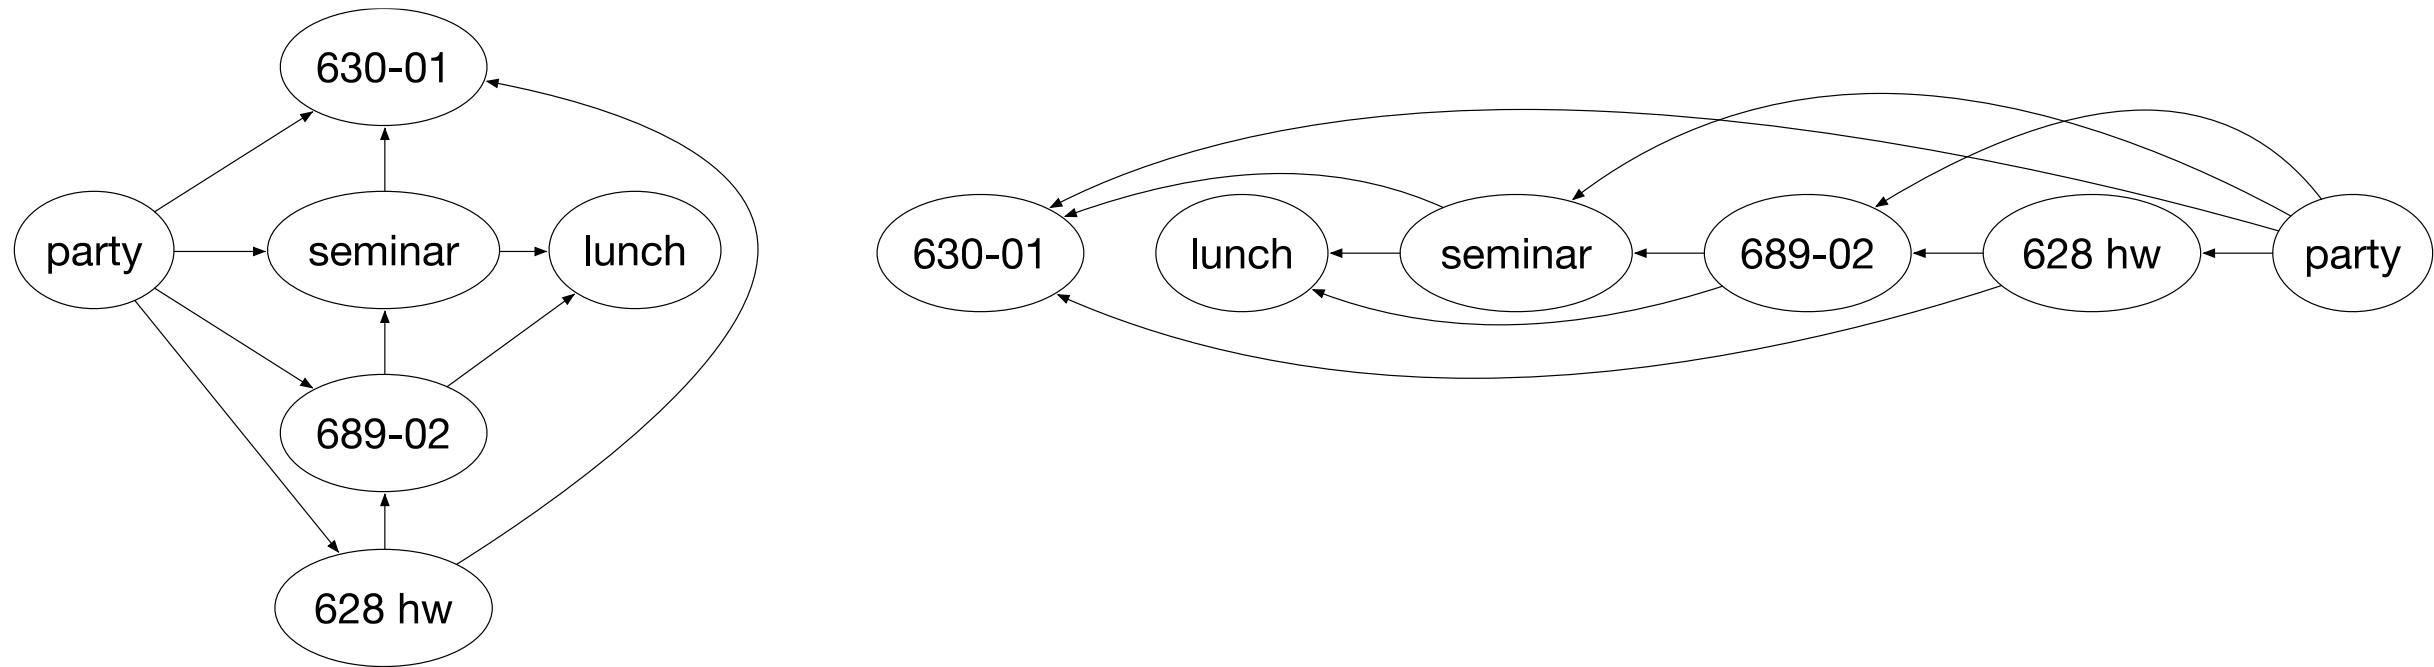
```
def shortest(root:Node, target:Node):  
    visited = {root}  
    worklist = [[root]] # list of paths  
    while len(worklist)>0:  
        path = worklist.pop(0)  
        p = path[-1] # tail of path  
        if p is target: return path  
        for q in p.edges:  
            if q not in visited:  
                worklist.append(path+[q])  
                visited.add(q)
```

Topological sort (acyclic graphs)

- **Problem:** Find linear ordering of nodes in directed acyclic graph such that all constraints, $u \rightarrow v$, are satisfied where u depends on v so v must come before u OR $u \rightarrow v$ mean u precedes v
- Examples: task ordering or course prereq chain.
- E.g., 502 is prereq for 621 and 601... Find order we should take classes
- Sort is not usually unique
- Cycles are meaningless for dependencies; how can 630-01 be attended before itself?



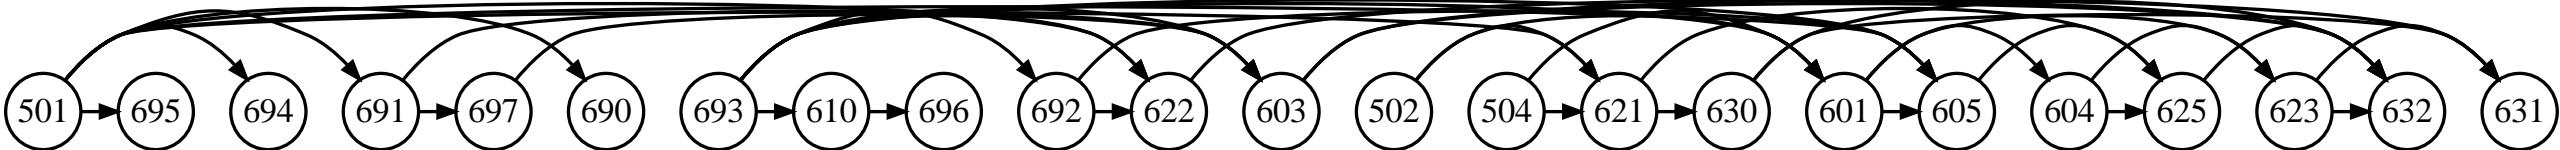
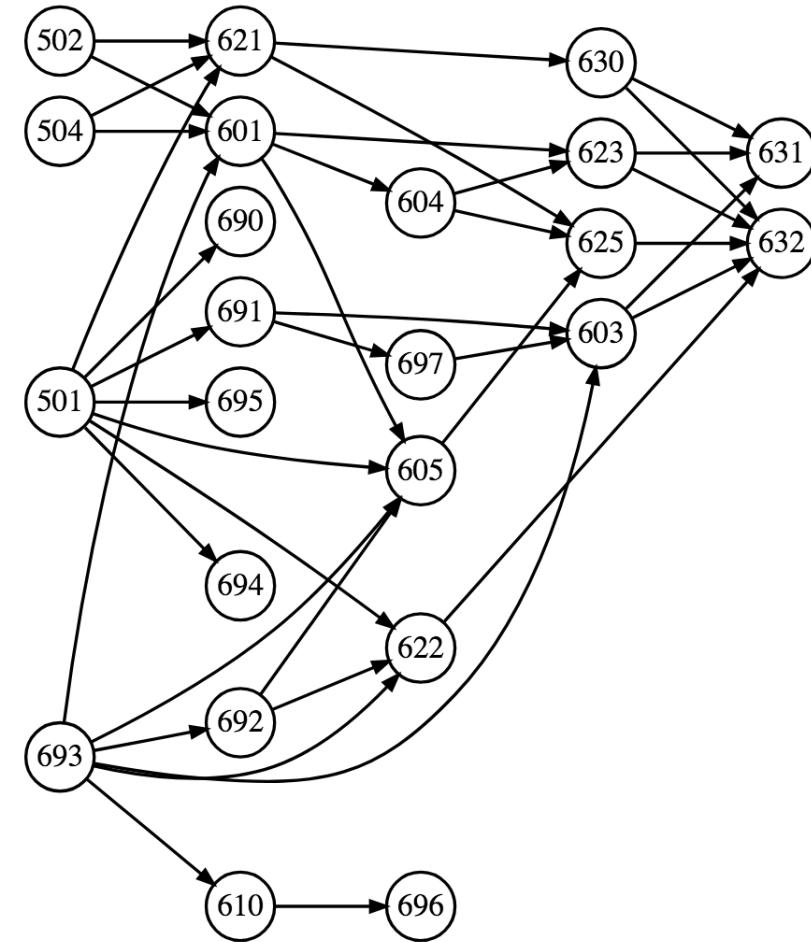
Example topological sort (u depends v)



If u depends on v, any linear ordering where edges point to left is solution

Example where u precedes v

- In this case, edges must point to the right



DFS postorder traversal

```
def postorder(p:Node, sorted:list, visited:set):
    if p in visited: return
    visited.add(p)
    for q in p.edges:
        postorder(q, sorted, visited)
    sorted.append(p) # postorder done after kids
```

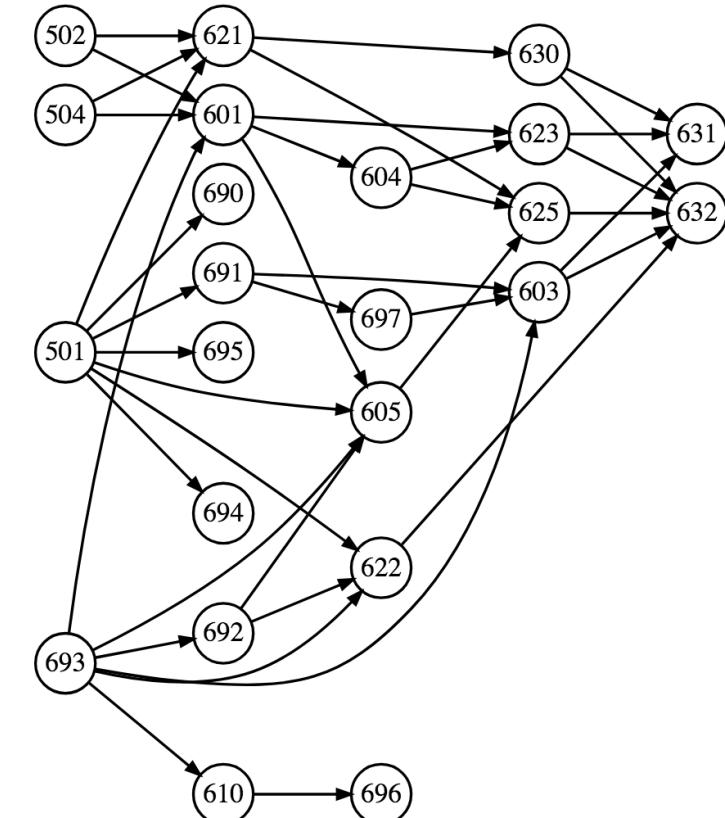
(recall: no cycles)

With multiple roots, hit them all

```
# nodes is dict edge list mapping from->to
def toposort(nodes:dict):
    sorted = []
    visited = set()
    while len(visited) < len(nodes):
        todo = [node for node in nodes.values()
                if node not in visited]
        if len(todo)>0:
            anode = todo[0] # pick a node
            postorder(anode, sorted, visited)
    return reverse(sorted)
```

We reverse postorder here since u precedes v

MSDS = "'''501->690
502,504,693->601
601->604
601,604->623
.....
..."



Summary

- Graphs are for showing relationships between elements
- DFS for finding a path or multiple paths or cycles
(recursive backtracking to find all nodes)
- BFS for find shortest (in edges) path or neighborhood
- DFS postorder great for topo sort
- Recursive alg's all use **visited** set or similar to avoid cycles
- Non-recursive DFS: (use work list stack)
 - push targets in reverse order onto work list
 - pop last work list item for next node to process
- Non-recursive BFS: (use work list queue)
 - push targets in order onto work list
 - pull from first position

Sample graph problems



Exercise

- Given a directed graph, detect all direct or indirect cycles
- For p in nodes:
report `iscyclic(p)`

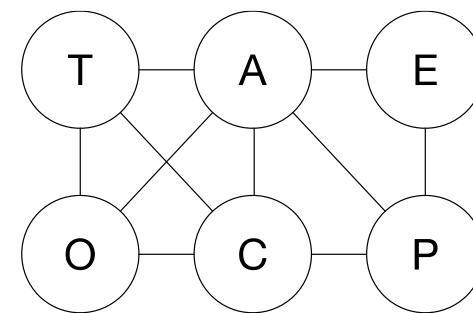
Exercise

- Given 2 lists P,Q and function $\text{conn}(p,q)=\text{true}$ if edge $p \rightarrow q$ exists. P is origin (starting) nodes and Q destination nodes. Report True for $P[i]$ reaches $Q[i]$ directly or indirectly
- E.g., P, Q could be cities and $\text{conn}(p,q)=\text{true}$ if direct flight $p \rightarrow q$
- Create graph using $\text{conn}(p,q)$ for all nodes in P and Q
- For each $P[i]$, see if $Q[i]$ is in $\text{reaches}(P[i])$ set.

Exercise: Boggle

- Given $m \times n$ matrix of letters. Find all English words possible by taking one adjacent step to another letter, starting with any letter; one occurrence of each letter per word; you're given a dictionary (/usr/share/dict/words)

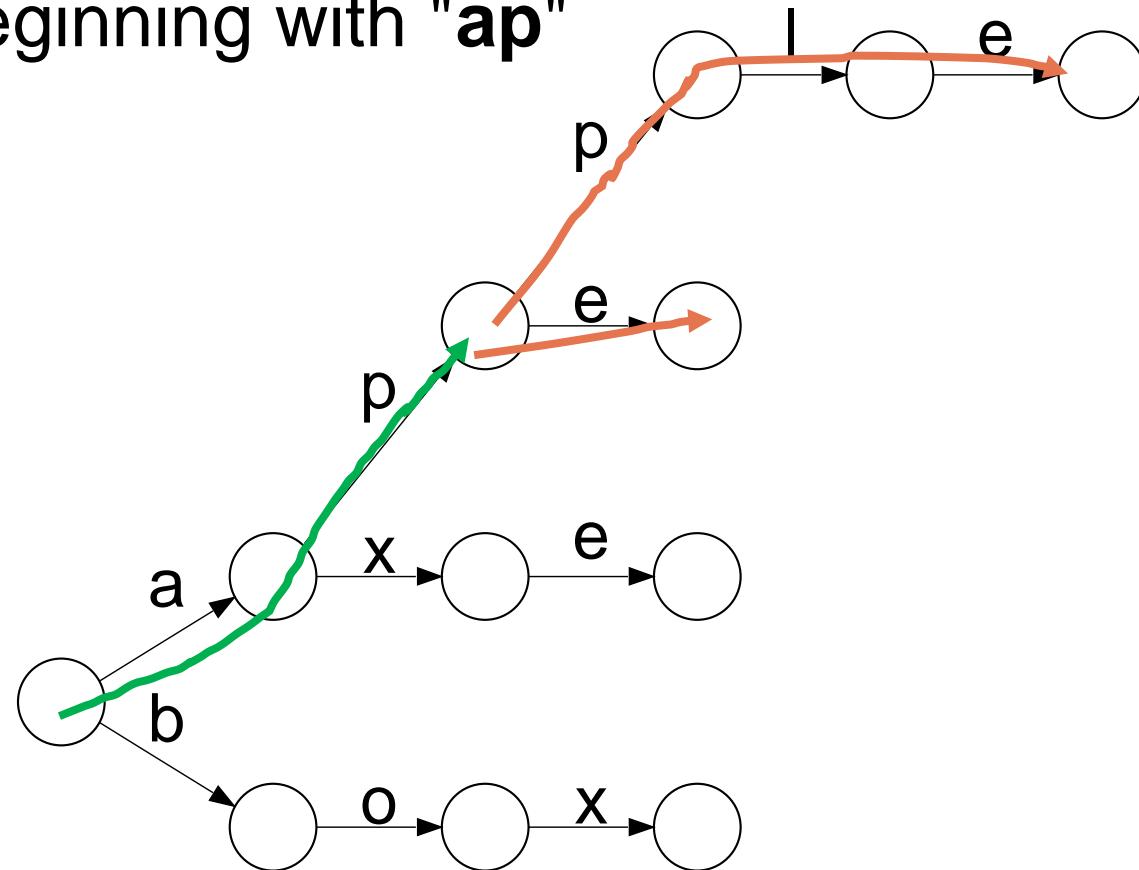
T	A	E
O	C	P



- For each node in graph, find all words
- For a specific starting node p, perform cycle-avoiding DFS; at each node q, look up word consisting of all letters on path from p to q, checking for duplicated letters

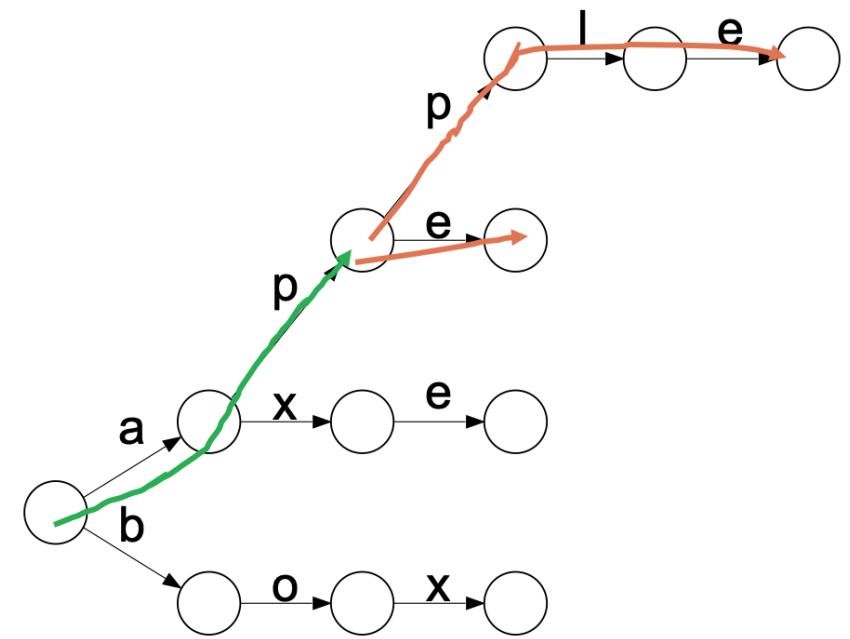
Exercise: *Find all words with given prefix*

- Find words beginning with "ap"



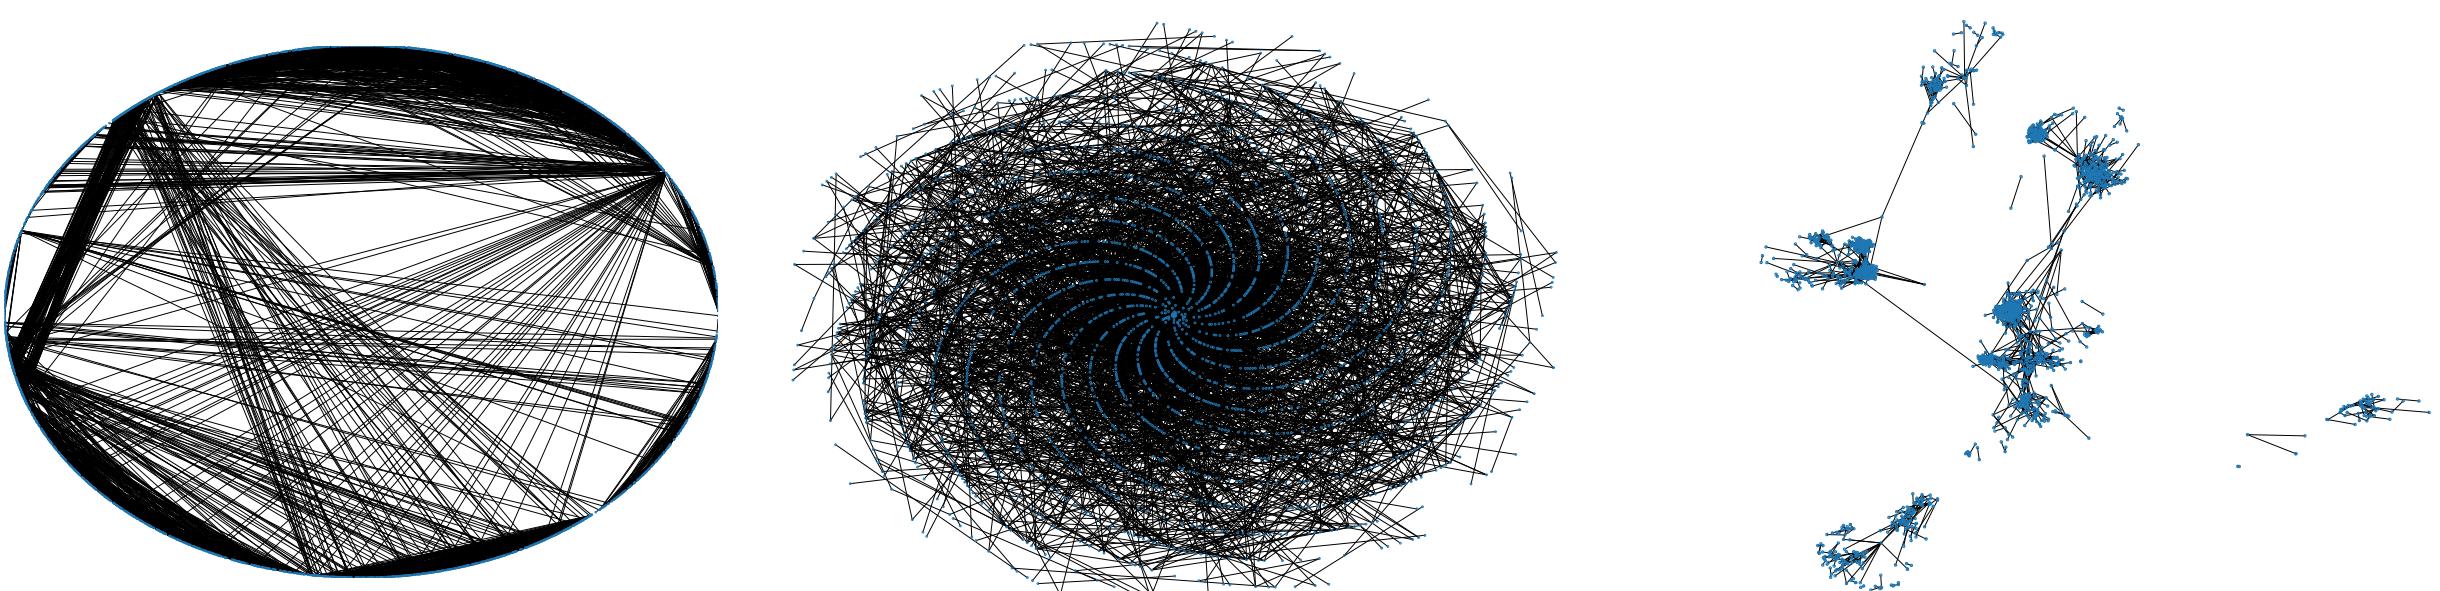
Exercise: Find all words with given prefix

- Build a Trie with dictionary words
- Find node in trie reached by prefix; call it p (green line)
- Recursively, using depth first search, find all reachable leaf (stop/accept) nodes (orange lines)
- Pass a path string as arg down the recursion chain to incrementally build strings reachable from p
`def suffixes_(p:TrieNode, path, paths): ...`
- There can be no cycle so we don't worry about "visited" arg, just track the current path and the overall list of paths we found



Coding exercise: Visualize: FB friend graph with networkx

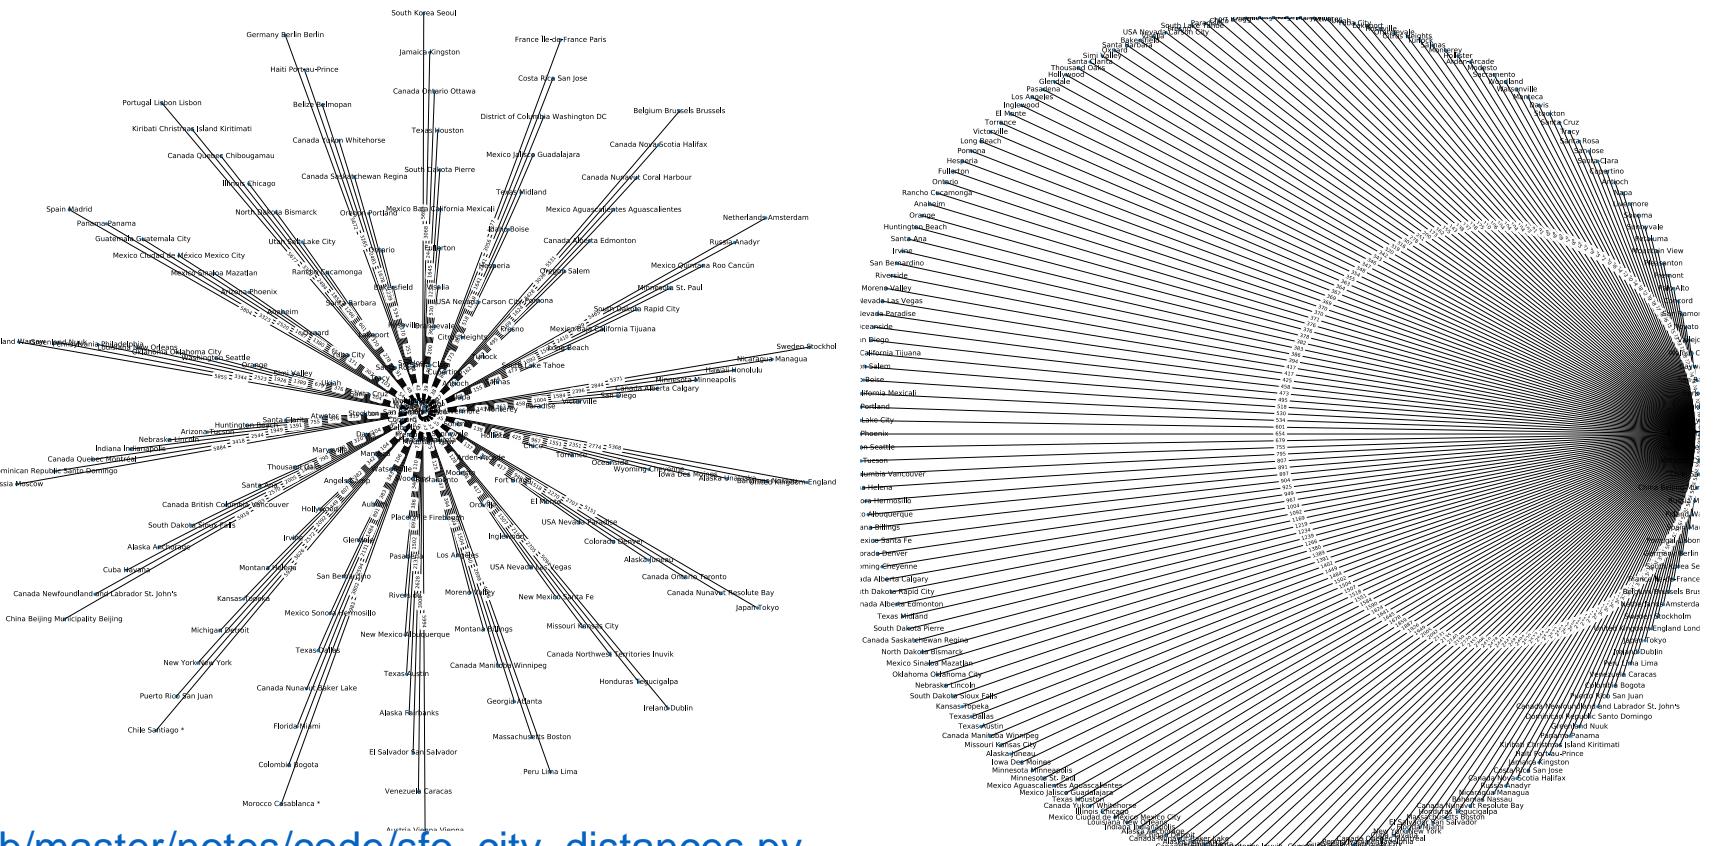
- pip install networkx
- download and uncompress https://snap.stanford.edu/data/facebook_combined.txt.gz
- Get small sample of edges then do edge_subgraph(), draw_networkx_edges(), draw_networkx_nodes()
- Need to pass positions of nodes for layout, such as circular_layout()



See https://github.com/parrt/msds689/blob/master/notes/code/viz_facebook.py

Coding exercise networkx: distances from SF to other cities

- Download data: <https://github.com/USFCA-MSDS/msds689/blob/master/notes/code/distances.csv>
- Create graph showing SF to other cities using multiple layouts



See https://github.com/parrt/msds689/blob/master/notes/code/sfo_city_distances.py